

Computer Graphics (0368-3236) Fall 2019, Exercise I

In this exercise you will implement basic image processing operations, such as scaling and rotation.

You will implement several sampling methods, such as Nearest Neighbour, Bilinear, and Parametric Gaussian.

As you've seen in class, it is preferred to use backward mapping instead of forward mapping, which will guarantee that the target image will be complete. You will display your results, to verify your code.

You will use Python 3 on a Jupyter Notebook environment with scientific packages. There are two ways to get started:

1. Get the **conda** package management and environment system conda from [here \(https://docs.conda.io/projects/conda/en/latest/user-guide/install/#regular-installation\)](https://docs.conda.io/projects/conda/en/latest/user-guide/install/#regular-installation). Conda is the preferred tool for data and algorithms development in the academia and the industry, but we can not provide support for students who have difficulties with it.
2. Alternatively, you can manually install the needed packages. We will use **Python 3.7** or above, and some very useful packages such as **OpenCV** or **Pillow** for image deserialization, **Imageio** for animation image serialization, **IPython** for notebook support, **Matplotlib** for plotting, **NumPy** for tensor data structures, **SciPy** for statistics functions, and **scikit-image** for image comparison metrics.

After installing python, run these commands:

- pip install pip --upgrade
 - or: python -m pip install pip --upgrade
 - or: python3 -m pip install pip --upgrade
 - or: pip3 install pip --upgrade
- pip install jupyter numpy scipy matplotlib opencv-python pillow scikit-image imageio --upgrade
 - or: python -m pip install jupyter numpy scipy matplotlib opencv-python pillow scikit-image imageio --upgrade
 - or: python3 -m pip install jupyter numpy scipy matplotlib opencv-python pillow scikit-image imageio --upgrade
 - or: pip3 install jupyter numpy scipy matplotlib opencv-python pillow scikit-image imageio --upgrade

Please post your questions on Moodle.

These are some imports, classes and functions implemented for you. You do not need to edit this section. Specifically, **you may not add imports anywhere else in the notebook.**

```

In [1]: # Copyright (c) 2019 Mattan Serry,
# Computer Graphics and Vision Lab, School of Computer Science, Tel Aviv University
#
# Permission is hereby granted, free of charge, to any person obtaining a copy
# of this software and associated documentation files (the "Software"), to deal
# in the Software without restriction, including without limitation the rights
# to use, copy, modify, merge, publish, distribute, sublicense, and/or sell
# copies of the Software, and to permit persons to whom the Software is
# furnished to do so, subject to the following conditions:
#
# The above copyright notice and this permission notice shall be included in all
# copies or substantial portions of the Software.
#
# THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR
# IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY,
# FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE
# AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER
# LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM,
# OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE
# SOFTWARE.

from collections import Counter
from math import sqrt, ceil, floor, radians, cos, sin
from os import remove
from sys import float_info, version_info
from random import randint, uniform
from time import process_time, sleep
from unittest import TestCase
assert version_info >= (3, 7)

# from cv2 import imread, cvtColor, COLOR_BGR2RGB
from IPython.display import HTML, display, Image
from matplotlib import rcParams
from matplotlib.pyplot import show, figure
from numpy import ndarray, iinfo, full, clip, average, array, empty, errstate, uint
8
from PIL import Image as PILImage
from scipy.stats import norm as normal_dist
from skimage import __version__ as skimage_version
assert skimage_version >= '0.16.2'

display(HTML("<style>.container { width:100% !important; }</style>"))

class Timer:
    def __enter__(self):
        self.start = process_time()
        print('*'*16)
        return self

    def __exit__(self, *args):
        self.end = process_time()
        self.interval = self.end - self.start
        print(f'Elapsed: {self.interval:.2f} seconds')
        print('*'*16)

class MyImage:

    def __init__(self, tensor: ndarray):
        self._tensor = tensor
        self._height, self._width, self._channels = tensor.shape
        self._time = None

```

Start by implementing a forward mapping operator. If you scale an image by **sy** and **sy**, where will each pixel translate to?

Part 1: Sampling Methods

You will need to implement the sampling methods *nearest_neighbour*, *bilinear* and *parametric_gaussian*. Make sure the returning value class is *PartialPixels*, which is an extension of *dict*. *PartialPixels* represents the weights of the source pixels per a target pixel, by method.

For example, NN should always return a *PartialPixels* of size 1, with weight 1.0.

Bilinear should return *PartialPixels* of size 1, 2 or 4, depends on the value of the source pixel. We relax the demand that the weights are summed to 1.

Parametric Gaussian should return *PartialPixels* of size up to 9, by calculating over a regular grid of 3x3 with stride **d**. Again we relax the demand that the weights are summed to 1.

We are providing some clipping and rounding methods for your assistance.

```
In [2]: def clip_to_range(z: float, *, max_dim: int):
        return clip(z, 0, max_dim - 1)

        def clip_and_round(z: float, *, max_dim: int):
            return clip_to_range(z, max_dim=max_dim).round().astype(int)
```

Implement *nearest_neighbour* and validate it. We recommend you use *clip_and_round*.

```

In [3]: def nearest_neighbour(height: int, width: int, y: float, x: float) -> PartialPixel
        s:
            # TODO
            return result

class NNTest(TestCase):
    def test(self):
        height, width = 768, 1024
        y1, x1 = -3.3, 58.12
        y2, x2 = 674.91, 3007.8

        nn1 = nearest_neighbour(height=height, width=width, y=y1, x=x1)
        nn2 = nearest_neighbour(height=height, width=width, y=y2, x=x2)

        print(nn1)
        print(nn2)

        # assert return type is class PartialPixels
        self.assertTrue(isinstance(nn1, PartialPixels))
        self.assertTrue(isinstance(nn2, PartialPixels))

        # assert that NN maps to exactly one pixel (the rounded), and the weight is
        # as expected
        self.assertTrue({(0, 58): 1.0} == nn1)
        self.assertTrue({(675, 1023): 1.0} == nn2)

NNTest().test()

{(0, 58): 1.0}
{(675, 1023): 1.0}

```

Implement *bilinear* and validate it. We recommend you to create a *Counter* instance, populate it, and then cast it to *PartialPixels*. We recommend you use *clip_to_range*.

```

In [4]: def bilinear(height: int, width: int, y: float, x: float) -> PartialPixels:
        # TODO
        return result

class BLTest(TestCase):
    def test(self):
        height, width = 768, 1024
        y1, x1 = 50.0, 60.0
        y2, x2 = 50.25, 60.0
        y3, x3 = 50.25, 60.50

        bl1 = bilinear(height=height, width=width, y=y1, x=x1)
        bl2 = bilinear(height=height, width=width, y=y2, x=x2)
        bl3 = bilinear(height=height, width=width, y=y3, x=x3)

        print(bl1)
        print(bl2)
        print(bl3)

        bl1.normalize()
        bl2.normalize()
        bl3.normalize()

        print(bl1)
        print(bl2)
        print(bl3)

        # assert return type is class PartialPixels
        self.assertTrue(isinstance(bl1, PartialPixels) and isinstance(bl2, PartialP
ixels) and isinstance(bl3, PartialPixels))

        # assert that NN maps to exactly one pixel (the rounded), and the weight is
        as expected
        self.assertTrue(
            {
                (50, 60): 1.0,
            } == bl1
        )

        self.assertTrue(
            {
                (50, 60): 0.75,
                (51, 60): 0.25,
            } == bl2
        )

        self.assertTrue(
            {
                (50, 60): 0.375,
                (51, 60): 0.125,
                (50, 61): 0.375,
                (51, 61): 0.125,
            } == bl3
        )

BLTest().test()

```

```

{(50, 60): 1.0}
{(50, 60): 0.75, (51, 60): 0.25}
{(50, 60): 0.375, (51, 60): 0.125, (50, 61): 0.375, (51, 61): 0.125}
{(50, 60): 1.0}
{(50, 60): 0.75, (51, 60): 0.25}
{(50, 60): 0.375, (51, 60): 0.125, (50, 61): 0.375, (51, 61): 0.125}

```

Implement *parametric_gaussian* and validate it. We recommend you to create a *Counter* instance, populate it, and then cast it to *PartialPixels*. We recommend you use *clip_to_range* for grid values, *normalnd* for the sampling value, and *bilinear* to project from fractional pixels to integer pixels.

Important: When you sample with *normalnd*, add **epsilon** to the result. Small values can cause numerical errors.

```

In [5]: def parametric_gaussian(height: int, width: int, y: float, x: float, std: float, d:
float, epsilon=float_info.epsilon):
    # TODO
    return result

from functools import partial
gaussian03 = partial(parametric_gaussian, std=0.3, d=0.025)
gaussian03.__name__ = 'gaussian 0.3'

class PGTest(TestCase):
    def test(self):
        height, width = 768, 1024
        y1, x1 = 0.001, 10.5

        pg = gaussian03(height=height, width=width, y=y1, x=x1)

        print(pg)

        pg.normalize()

        print(pg)

        # assert return type is class PartialPixels
        self.assertTrue(isinstance(pg, PartialPixels))

        # Verify values of the Gaussian sampling
        self.assertAlmostEqual(pg[(0, 10)], 0.5, places=2)
        self.assertAlmostEqual(pg[(0, 11)], 0.5, places=2)
        self.assertGreater(pg[(0, 11)], pg[(0, 10)])

        self.assertAlmostEqual(pg[(1, 10)], 0.0, places=2)
        self.assertAlmostEqual(pg[(1, 11)], 0.0, places=2)
        self.assertAlmostEqual(pg[(1, 10)], pg[(1, 11)])

PGTest().test()

{(0, 10): 7.858954823217127, (0, 11): 7.858954823217129, (1, 10): 0.071215723910
53246, (1, 11): 0.07121572391053246}
{(0, 10): 0.4955098239383813, (0, 11): 0.4955098239383814, (1, 10): 0.0044901760
61618693, (1, 11): 0.004490176061618693}

```

```
In [6]: def calculate_all_partial_pixels_by_method(height: int, width: int, method: callable, backward_map: dict) -> ndarray:
        result = {}

        for (new_y, new_x), (old_y, old_x) in backward_map.items():
            result[(new_y, new_x)] = method(height, width, old_y, old_x)

        return result
```

Part 2: Scale Transformation

We will experiment with a photo of a mandrill. Make sure you can load the image and see the monkey.

```
In [7]: from base64 import b64decode
        from io import BytesIO
```

```
mandrill = "/9j/4RxxRXXhpZgAATU0AKgAAAAgABwESAAMAAAABAAEAAAEaAAUAAAABAAAAYgEbAAUAAAA
BAAAAagEoAAMAAAABAAIAAAExAAIAAAAIAAAACgEyAAIAAAAAUAAAAIdpAAQAAAAABAAAAqAAAAANQACvyAAA
AnEAAK/IAAACcQQWRvYmUgUGhvdG9zaG9wIENDIDlwMTUgKFdpbmRvd3MpADlwMTk6MTE6MDggMDA6MjE6M
jEAAAQOgAQADAAAAAQABAACgAgAEAAAAAQAAAJugAwAEAAAAAQAAANAAAAAABBgEDAAMAAAABAAyAAAEa
AAUAAAABAAAABIGebAAUAAAABAAAABKgeOAAAMAAAABAAIAAAIBAAQAAAAABAAAABMgICAQAQAAAAABAAAFwAAAA
AAABIAAAAAQAAAEgAAAAB/9j/7QAMQWRvYmVfQ00AAf/uAA5BZG9iZQBkgAAAAAH/2wCEAAwICAgJCawJCQ
wRCwoLERUPDAwPFRgTExUTEExgRDAwMDAwMEQwMDAwMDAwMDAwMDAwMDAwMDAwMDAwMDAwMDAwMDAwBDQsLDQ4NE
A4OEBQODg4UFA4ODg4UEQwMDAwMEREMDAwMDAwRDAwMDAwMDAwMDAwMDAwMDAwMDAwMDAwMDAwMDAwMDP/AABEI
AKAAdwMBIGACEQEDEQH/3QAEAAj/xAE/AAABBBQEBABQEAQAAAAAADAACBAUGBwgJCgsBAAEFABQEBABQEB
BAAAAAABAAEAAgMEBQYHCAkKCxAAAAQQAABwIEAgUHBggFAwwzAQACEQMEIRIxBUFRYRMicYEyBhSRobFCIy
QVUUsFiMzRygtFDBYWSU/Dh8WNzNRaisoMmRJNUZEXCo3Q2F9JV4mXys4TD03Xj80Yn1KSftJXE1OT0pbXF1
eX1VmZ2hpamtsbW5vY3R1dnd4eXp7fH1+f3EQACAgECBAQDBAUGBwcGBTUBAAIRayExEgRBuWfXtHmFMoGR
FKGxQiPBUThWMyRi4XKCKkNTFWNzNPelBhaisoMHJjXC0kSTVKMXZEvvNnR14vKzhMPTdePzRpSkhbsVxNT
k9KW1xdXl9VZmdoaWprbGlub2JzdHV2d3h5ent8f/2gAMAwEAAhEDEQA/AOS+z1hi15Lgfa0ckj91a9Nm1z
WzILCdToqt1Qpt2Ws3Vu1Y/wAf/JKvSG336FzQ3VpGgiOFU4QiWPone91Rf1lbtS/SEka/99QrrrwQK3720
1mdY/rk9iiaurXU6H1CA555gmNrv5Sh1YVLANS1SQS4Tz3bwgThecdao8NjrXbXzPGp8FedUxtbmYQ0sQJ
Ht/FRLRQAGuLjETqSf6v8rak2Q4yXAgAbmhs6+Drvbu/ca90jiiBclDH1Q/Zdd1JLQ36Q8D3UjW5j97vc8C
RE6RrMKVYG/adoLYaJkxP5u1v/VfvpzZ+1cSd4AIBmTHGn8n8+t6RxxwkdqKpY92u7Ide9oktfzHkD/moxsc
0kv1AALSD4f+TT2YvqbXshrtA4jiJ921QyqHPAh20T7h8OyryhKQQR1MZMUQYyru28W7Ez6n42Xy4QxxOoP
5rm/wAtr1Uf0xuHWWfPpHAWLPEDX/Pwf1C1drW1aQ0EmO60fXvy6a6+X6EE/D81TS+Xt/BklVdiidlCj02
2Uklume8/HcrGT9ncxuxpY3byJn8EUmuuposZueNSDqZCpnOt9QmsAtPG7ny/qtUX6O6wfl+L1Qe1tZq9R7i
72z7pGnKsMLbTMsalw4Hy3f8ASSRv/GVxeOr/AP/QqCn7ZjMZkBrLGwAfyroAbVKNdYLD38D3LXmts9Z7ne
9pkR2j93RQyMJ8+u2WhzgHaaidGuUAiynytrA+viPFA3DfLB4wfaFodLaHUmmyJeJYTPb2ln735qj0/pjqL
JMEkSdvE/ym/nbv8IrxoFLHOAJgaM/8glIro+CLPRYur5Yrc8tEO0fEyOQD7fosY/8A6azsfqIs1Al75B7k
wfd7Xyx30Fof+dqfWG218Onse08Tp9H9LtKxWgstc08STESNPdx8k4DiViGbOxV1AUWWGY2yGtHtAJb7v3v
3WfRT0ZocCSIB01usiY3fnN+isJrzZL+J1IA/s/8AmK08Wix0QwughoA7uP8AJcgYgIt2KbfUlw0a3x1n+T
/UVioUuH6Tnlzo8/8AvqBhs9MnKbXiB2Gvt9hb9H/yC0PRYxulrWt9Ruh4iFDZMT6ZR6MWSylrOnV5OTvrE
VP7iCf6v8pUcjEyMTN3h2gJA29o7O3fyVbNmU0uaw+m36IESSP3m/yU3qOdXBh1jdPET+65VuKRHj+TGBId
WtTbk2P3kaTABmUfMxQ2sXVTG9w9wnnwRDRa6oWAwfydGkT+6nqx/Tta660ubOrSQZMTt/8AJpxPT8Uk6NN
jLLsV7HuLLq4cxwGpI92z+0ktI47BbLCPsIBd9/0kkaFVet/inijXE//Rz8C/Mra2W7iNYOv9SV0FbBm1hu
SxhBE7ODP5zXj/AMyQ8HHf7iHetu9opcCC0nlm2PfsRLq20UQBE+2OWx/I3fQ/qgPg0tfxXotVTRUD6HsIn
azB0A3+3d/JU3s04A/RaJaCHDR3Zrnf5xu9r/wCXs/wiGLTt1EyQ6TqQXDa72gn6X0FTzOobAYBG3XUCr3bt
3N/N/wCDfvs0pBLm9Z9Evd4iGjQj91vub/0vTXJ5ZAc+By0sA83nstTqWTdc4kDThvz76fyVl24pex2/kG
T930QpIRJGgyRCfptTbNu6I/OmDxq0Lrem4FdVQc5sHUN4D6Ubd30foMd/wCCLlOmC2mxu76Mwef4Ls8PLY
a2NLZDnAdwSf5G4/u/mf8AnCYd9UshjD2BrZdpE8cT7nfR771HaLAWatc0kt2kawI9mQxzNrdCQI0kAcgg
fnV+33qlc8bzs0giJM+befpf+CVv/nHp9RMTE6grNUfUGfoK76/eWuDXR3B0c3s72oTultNe5ziysjcW8Qt
fpoORS4OJG4y0n9wRLpb9Hcg5nTsiu000lYglhGpn+W5Uz8vISJA0RWq2L099GK47i+u3SewlCtwgMiqwG
wscP87/yTvc03WZdNTgxw+ziA5oJIAp3GH85yEy37TW21lhAYIdX2lV8plRrTZrj1BrvbbvvpbdXitbAe0A
2E8xodv8lv7ySy8j1mZNOM502004Edmn6M/wBZJNolayvKn//SvMyW1gB28AeXuA/e/qrRrgMqlsndJOwx3
j935+9ZuORZuc33FpkSDJPh6h/6pqu1j08Wtto9wBe9sGJCpO/2GpspekDum1ltGAxjWtMtMmBAJ2kfPng/
Tc57v9fesyykvc4tHA0brwNZ2/9+Rep5Wyt8EHaNQRP0Z2/9+ZXW7/hVz1XlOwaR+ssc+yfdiJd/a3e121
GA8aWybz8azJ3HHqn90mB/V0QndJy2taxlRa7lxIGnxCu90+u31b9Mtubdj3QYt0e3tDRG13u/fej5P15+r
BYNzr7nbdGVsa0f1f0hU0eEDdbwE62Gng0YtZit9jmEbXEdvj7l0U47Xels1kkQBjJP8AgzPt/s3f9bWU/
wCt31bs1ZRelxHt9Qtir+c81+76X8pQo6o23KFja341TvolzXsYQT9Cqyx+immAA0IKLN6indLCXENPuPc
agu0d00h27a3/P8A0jFWur3NhpGGS1w9zRInXX/z3+Yilu3iBq46Rtnvv0Yzd+fv/wDRPp/QT2NsL9r2PLy
NK3sewKSGy2rayzZve/8AqJlaJttDLbbuDahpM6GX13Ie+PZ7Gu/Rf4FXMmpzZlZs7jgkxzMu0cxu5UMK80
3FlgIeBOx52d/c38123/v6v3Xh7SZ2zIluoI/s/wDf1JvWernFrTYaLA42uA9oaOCY/N+isa/1KH2VMZNY
B08dvf/ADlvOxtobaCZ1AaHTHwfDd2//oLKuxHsZ6wPuJJCp3hPI/te1v76zJw9ZvquABNFz2dQY/G+0tEW
t3BxIG46T/1SStHHoaQYDbHauG4abv5LP6ySjqNdd6pVrt//09WjCLWB1gLQ7T7/AG+0mfTY5v8AbQOoXel
LQOOA4SZGu8tn+T6q2RU54ddYPa0e0bvotPw9nuWD1d8F5Y6Hg69/+p/cdsTJdAuGxcPqDiQ4awJiOB+7x/
n/APXP+MXP0dCz+tZLmdMxX5Lw4tNjQ1tQcPdsdkWluO3/Aiv1Peul6J0Wz6xdUdgGx1OLQ0W5TgTJrmPQp
/lX/Q3P/ml6NViY1OPj049TKaWF7WVVtDWTj3Mc1g/OZ+99NTYoXqdmGciNt3xbq31b61h5ONZ1KttdmS8V
UmpzXzYCP0VkoLmP93/Fo2V9UsvI+teX0jpdQssxXNe4WOHptaA11z8i327Kmuds/wCoXovU6a+qfXvpWBs
Bq6WHZ2S48CwhtrP7Xto+kjfvJGDmdR61t3WdZy7Lcdp/0FRdXTY/d/L9S3zez+bSAiZ0Bo2JQl1j5fikQZy
4JbfLx8df8x846r9U+udIuHU8jHYKanMcy+gtsqY4OGyzJG1jtm7/gPTWx0762fWSow1/bBdTc0gnIBtLTq
XFr3Ncx/wDVdUu56jZj1VvbfFrLWltpd7g40EOZDvdZv+ivKMy6vB61biYj3WV0vH2ciCZ+1XW5rva59T/Z
YpMuERiCoRuxZpGRiasPTY/1qzMINDn90qNZJf62NAFulvev9Bd6jP5awrevftW77Vfm3UZwYa2F0vYGA7h
T/pmez+c9Kzeq3VH5uRmlMyr6si5lWyRu27iS/wBJtr/c+ze76X7/APNqjXhUmXWAVOaYoyYjSQ5pUBMQGb
ht6bBz+p07t2Yy1xlzK64cySPe+b/Vs9R7G/mf+fFo9N6q3Nc/GzIrvDg6r91wd9KrsNu13u97FxnPp9tU
WNomzuB5s9rdqLiZrCA8keozh7ZBjh0N9rtr1EZYuwVERVF7gteS7GBAeNDu8vcGtHt9jvc/wBr1rvzIeca
wkREOLsa16LRWad4PVG5dYzeY1Z0TRFad6N92i.TmYF9wORU5ocwROR7cfnP3f6iaZY3PishinRi+xcHHTst+0
```




Implement the *transform* function:

1. Receive an image object and a backward map (a dictionary that maps pairs (y, x) to partial pixels)
2. Create a new empty numpy tensor in a size according to the keys of the backward map
3. Decide the color for every coordinate, by calling *get_average_color* with the partial pixels and the source image
4. Return a new image object

```
In [8]: def transform(source_image: MyImage, backward_map: dict) -> MyImage:
        # TODO
        return target_image
```

Implement a simple forward scale mapping. Scale **x** by **sx** and scale **y** by **sy**.

```
In [9]: def forward_scale_mapping(old_y: float, old_x: float, sy: float, sx: float) -> (float, float):
        # TODO
        return new_y, new_x
```

Now implement a backward mapping operator. Given a target image and known scales, where did each pixel translated from? The result may be a "fractional" pixel.

You can implement this function by a single call to *forward_scale_mapping*, think how. Partial points will be given to students who do not call *forward_scale_mapping*.

```
In [10]: def backward_scale_mapping(old_y: float, old_x: float, sy: float, sx: float) -> (float, float):
        pass # TODO
```

We will now test your code for sanity.

```
In [11]: class ScaleTest(TestCase):
        def test(self):
            y, x = randint(0, 768 - 1), randint(0, 1024 - 1)
            sy, sx = uniform(0.1, 10), uniform(0.1, 10)

            new_y, new_x = forward_scale_mapping(y, x, sy, sx)
            old_new_y, old_new_x = backward_scale_mapping(new_y, new_x, sy, sx)
            self.assertAlmostEqual(y, old_new_y, msg="Diff in y")
            self.assertAlmostEqual(x, old_new_x, msg="Diff in x")

ScaleTest().test()
```

Now let's create a full backward mapping dictionary. Given a source image's dimensions, you need to calculate the target image's dimensions, and then for every target pixel (**new_x**, **new_y**), store the source pixel (**old_x**, **old_y**) in the dictionary.

```
In [12]: def calculate_all_backward_scale_mapping(source_height: int, source_width: int, sy:
float, sx: float) -> dict:
    """

    :rtype: (dict, int, int)
    :param source_height: Original height of the image in pixels
    :param source_width: Original width of the image in pixels
    :param sy: Forward height resize scale
    :param sx: Forward width resize scale
    :return: A mapping from forward integer pixels (y, x) to the backward fractional
    pixels (old_y, old_x)
    """

    # TODO

    return backward_map
```

Let's scale the mandrill using 3 sampling methods.

```
In [13]: sy, sx = 0.9, 2.1

run(mandrill)

old_height, old_width = mandrill.height, mandrill.width
backward_map = calculate_all_backward_scale_mapping(old_height, old_width, sy, sx)

run(mandrill, nearest_neighbour, backward_map)
run(mandrill, bilinear, backward_map)
run(mandrill, gaussian03, backward_map)
```



PSNR = inf dB
 MSE = 0.00
 Elapsed: 0.11 seconds

 nearest_neighbour



Elapsed: 5.19 seconds

 bilinear



Elapsed: 5.53 seconds

 gaussian 0.3



Elapsed: 196.56 seconds

Part 3: Rotation Transformation

Implement a simple forward scale mapping. Given **x** and **y**, translate them (use the **height center** and **width center** as the origin) and rotate them using the **cosine of t** and the **sine of t**.

```
In [14]: def forward_rotation_mapping(old_y: float, old_x: float, height_center: float, width_center: float, cost: float, sint: float) -> (float, float):
          # TODO
          return new_y, new_x
```

Now implement a backward mapping operator. Given everything, where did each pixel translated from? The result may be a "fractional" pixel.

You can implement this function by a single call to *forward_rotation_mapping*, think how. Partial points will be given to students who do not call *forward_scale_mapping*.

```
In [15]: def backward_rotation_mapping(new_y: float, new_x: float, height_center: float, width_center: float, cost: float, sint: float) -> (float, float):
          pass # TODO
```

We will now test your code for sanity.

```
In [16]: class RotationTest(TestCase):
    def test(self):
        height, width = randint(1, 768), randint(1, 1024)
        y, x = randint(0, height - 1), randint(0, width - 1)
        t = uniform(0, 360)

        cost, sint = cos(t), sin(t)

        height_center = (height - 1) / 2
        width_center = (width - 1) / 2

        new_y, new_x = forward_rotation_mapping(y, x, height_center, width_center,
        cost, sint)
        old_new_y, old_new_x = backward_rotation_mapping(new_y, new_x, height_center,
        width_center, cost, sint)
        self.assertEqual(y, old_new_y, msg="Diff in y")
        self.assertEqual(x, old_new_x, msg="Diff in x")

RotationTest().test()
```

Now let's create a full backward mapping dictionary. Given a source image's dimensions, you need to calculate the target image's dimensions, and then for every target pixel (**new_x, new_y**), store the source pixel (**old_x, old_y**) in the dictionary.

```
In [17]: def calculate_all_backward_rotation_mapping(height: int, width: int, t: float) -> dict:
    """
    :rtype: dict
    :param height: Height of the image in pixels
    :param width: Width of the image in pixels
    :param t: Forward rotation angle in degrees (in relation to the center of the image)
    :return: A mapping from forward integer pixels (y, x) to the backward fractional pixels (old_y, old_x)
    """
    # TODO

    return result
```

Let's begin rotating the mandrill 10 times, 36° each, with our sampling methods. We expect to see unwanted artifacts near the edges, because the image is not padded.

```
In [18]: t = 36
assert 360 % t == 0
iters = 360 // t

run(mandrill)

backward_map = calculate_all_backward_rotation_mapping(mandrill.height, mandrill.width, t)

run(mandrill, nearest_neighbour, backward_map, num_iters=iters, gif=True)
run(mandrill, bilinear, backward_map, num_iters=iters, gif=True)
run(mandrill, gaussian03, backward_map, num_iters=iters, gif=True)
```



```
PSNR = inf dB
MSE = 0.00
Elapsed: 0.08 seconds
*****
*****
nearest_neighbour
```

```
<IPython.core.display.Image object>
```

```
PSNR = 16.21 dB
MSE = 1557.86
Elapsed: 16.86 seconds
*****
*****
bilinear
```

```
<IPython.core.display.Image object>
```

```
PSNR = 16.76 dB
MSE = 1371.52
Elapsed: 18.77 seconds
*****
*****
gaussian 0.3
```

```
<IPython.core.display.Image object>
```

```
PSNR = 16.76 dB
MSE = 1371.72
Elapsed: 130.53 seconds
*****
```

Implement the ***pad*** function. It should take an image and return a larger image, having the original image in the center and white pixels around it, such that no pixels will be lost during a rotation operation.

1. Hint: Where would a corner pixel translate under a 45 degrees rotation? Think of a square source image for example.
2. Hint: Should the target image be square? Is there symmetry in the problem?

Partial points will be given for a larger padding than needed.

```
In [19]: def pad(image: MyImage) -> MyImage:
          # TODO
          return MyImage(new_tensor)
```



```
In [20]: padded_mandrill = pad(mandrill)

run(padded_mandrill)

backward_map = calculate_all_backward_rotation_mapping(padded_mandrill.height, padded_mandrill.width, t)

run(padded_mandrill, nearest_neighbour, backward_map, num_iters=iters, gif=True)
run(padded_mandrill, bilinear, backward_map, num_iters=iters, gif=True)
run(padded_mandrill, gaussian03, backward_map, num_iters=iters, gif=True)

*****
```



```
PSNR = inf dB
MSE = 0.00
Elapsed: 0.08 seconds
*****
*****
nearest_neighbour

<IPython.core.display.Image object>

PSNR = 164.01 dB
MSE = 183.12
Elapsed: 43.55 seconds
*****
*****
bilinear

<IPython.core.display.Image object>

PSNR = 165.10 dB
MSE = 142.63
Elapsed: 50.56 seconds
*****
*****
gaussian 0.3

<IPython.core.display.Image object>

PSNR = 164.89 dB
MSE = 149.70
Elapsed: 354.92 seconds
*****
```

Bonus: load an image of your choice and define your own Gaussian, such that for 10 rotations, the PSNR is higher and the MSE is lower in comparison to bilinear. Attach the image to your submission.

Partial points will be given for unsuccessful attempts.

```
In [21]: # The Gaussian's standard deviation
my_std = 1.0
# The distance of the regular grid's sampling
my_d = 1.0
# The epsilon to add to the sampling result
my_epsilon = 1.0

my_gaussian = partial(parametric_gaussian, std=my_std, d=my_d, epsilon=my_epsilon)
my_gaussian.__name__ = 'my gaussian'

my_image = PILImage.open("path/to/your/image.jpg").convert("RGB")
my_image = MyImage(array(my_image))
my_image = pad(my_image)

run(my_image)

backward_map = calculate_all_backward_rotation_mapping(my_image.height, my_image.width, t)

run(my_image, bilinear, backward_map, num_iters=iters, gif=True)
run(my_image, my_gaussian, backward_map, num_iters=iters, gif=True)

-----
FileNotFoundError                                Traceback (most recent call last)
<ipython-input-21-83f8fd531c8c> in <module>
      9 my_gaussian.__name__ = 'my gaussian'
     10
--> 11 my_image = PILImage.open("path/to/your/image.jpg").convert("RGB")
     12 my_image = MyImage(array(my_image))
     13 my_image = pad(my_image)

~\Anaconda3\envs\python37new\lib\site-packages\PIL\Image.py in open(fp, mode)
    2764
    2765     if filename:
-> 2766         fp = builtins.open(filename, "rb")
    2767         exclusive_fp = True
    2768

FileNotFoundError: [Errno 2] No such file or directory: 'path/to/your/image.jpg'
```